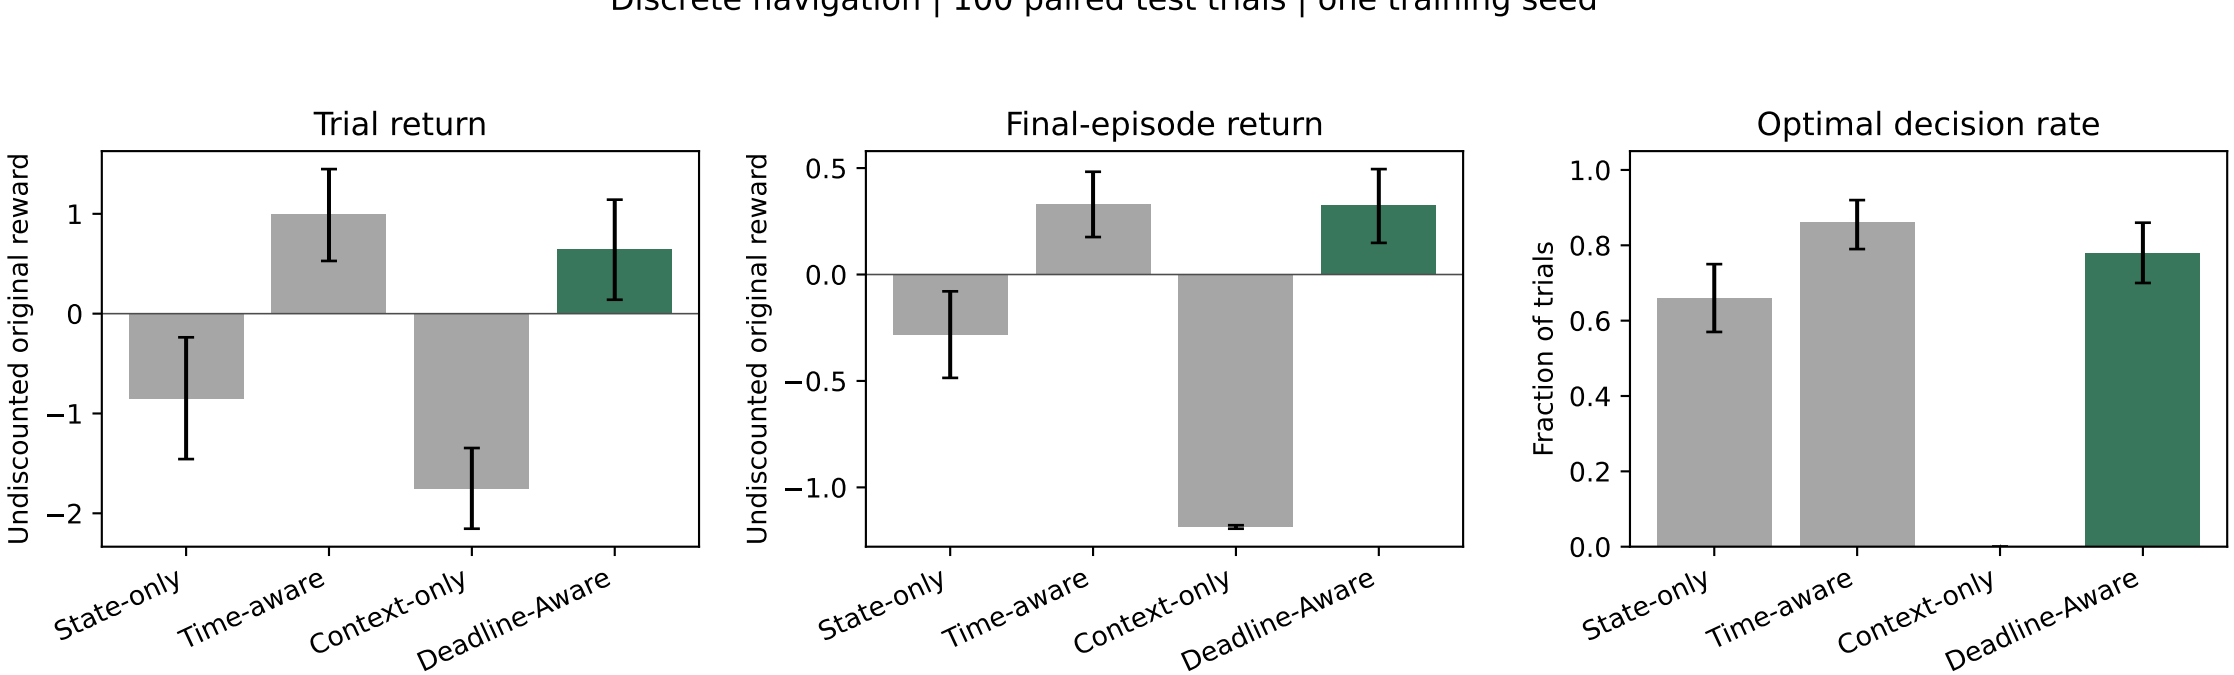
Discrete navigation | 100 paired test trials | one training seed
Trial return
Undiscounted original reward
1
0
−1
−2
State-only
Time-aware
Context-only
Deadline-Aware
Final-episode return
Undiscounted original reward
0.5
0.0
−0.5
−1.0
State-only
Time-aware
Context-only
Deadline-Aware
Optimal decision rate
Fraction of trials
1.0
0.8
0.6
0.4
0.2
0.0
State-only
Time-aware
Context-only
Deadline-Aware
Error bars: 95% trial-bootstrap CI. Returns retain step costs; ODR uses time-neutral route optimality.