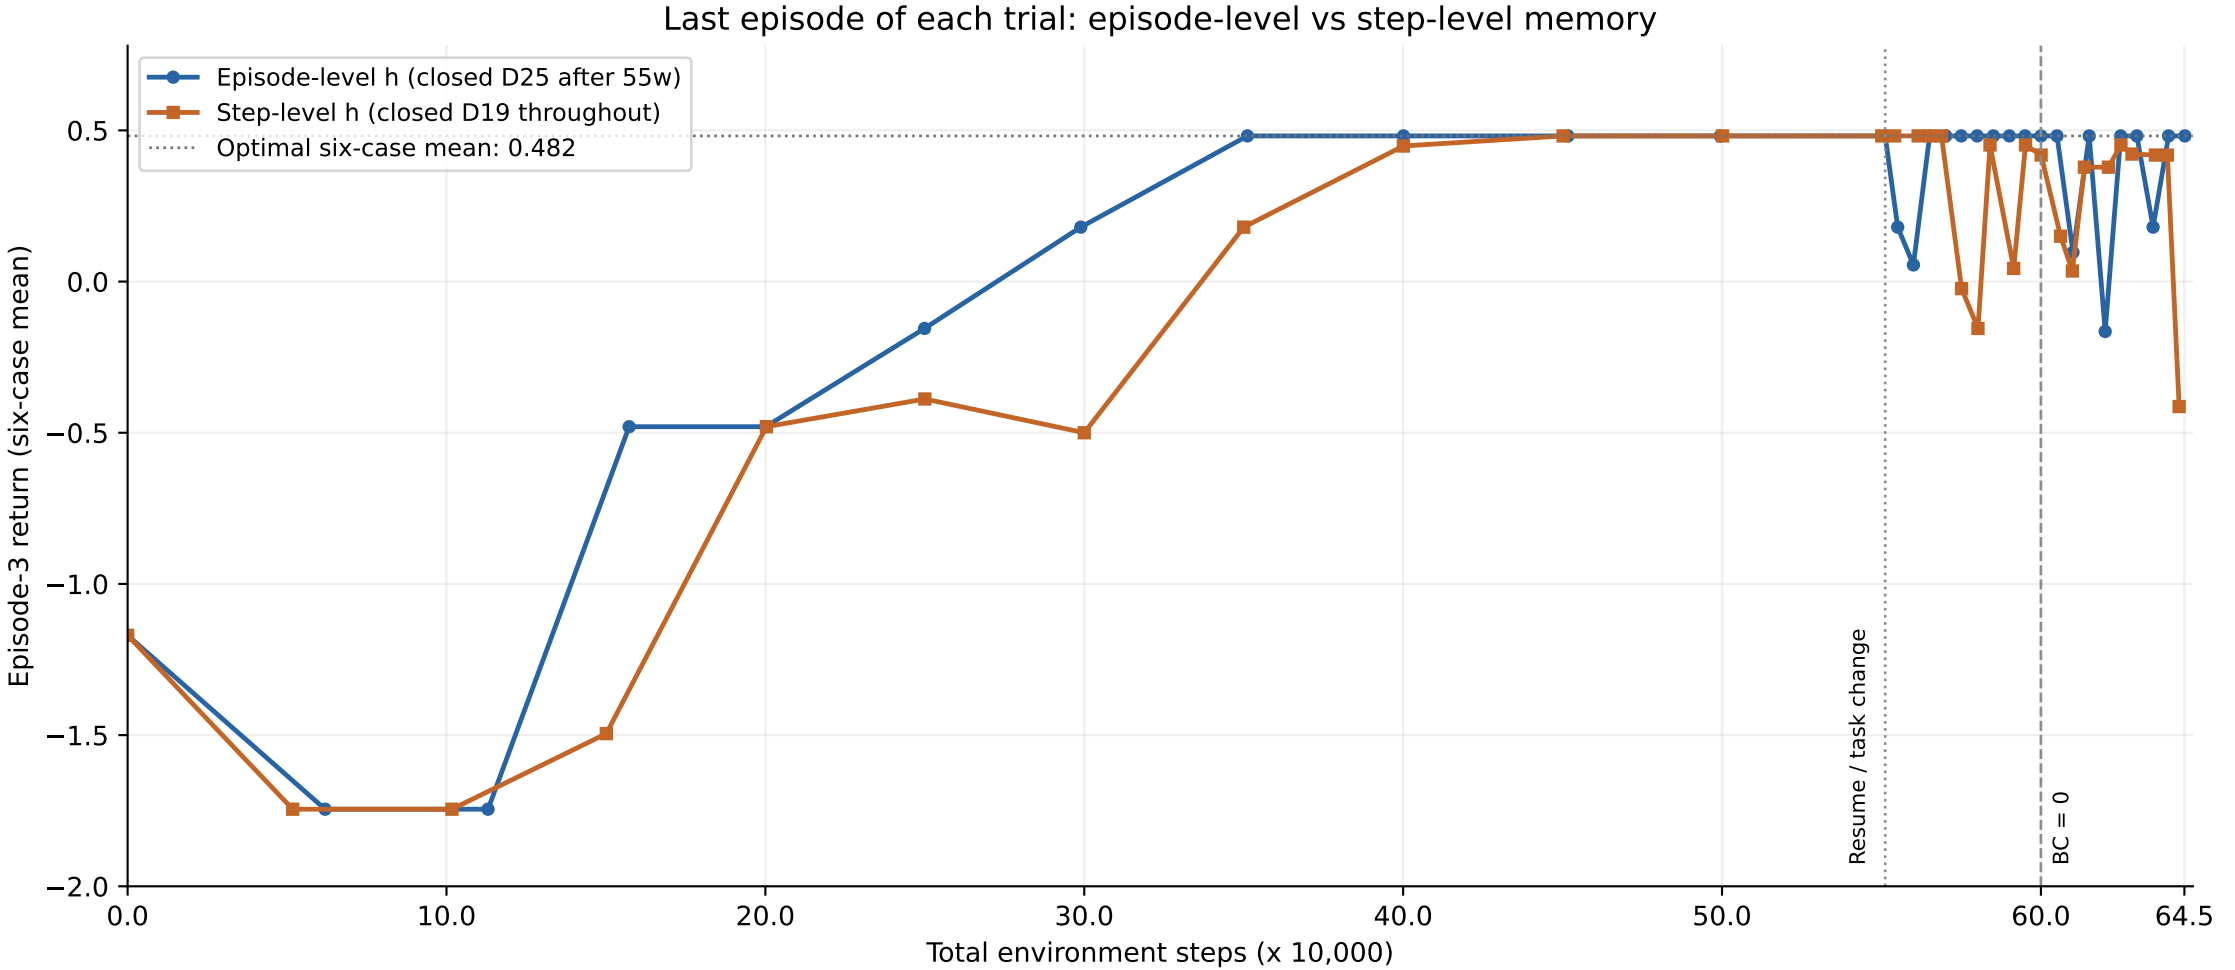

Source: original fixed-task evaluations + the latest episode-level D25 branch + saved step-level D19 weights.
Only episode 3; 6 cases equally weighted; 3 evaluation seeds/case; deterministic actions; no smoothing.
Tasks differ after 55w (closed D25 vs D19): this is NOT a matched-task architecture ablation.
Zero-step networks reconstructed with seed 0. Markers are actual checkpoints; no extrapolated endpoints.